

1. Host and target setup	1h26m	
1. About the instructor		
2. Source code and course materials		youtube
3. Host and target setup		1
4. Tool-chain download		2
https://releases.linaro.org/components/toolchain/binaries/7.5-2019.12/		
5. Important Note		
6. Installing gparted application		
7. Tool-chain installation and PATH settings		3
8. Note for the students		
9. Target preparation : Serial debug setup		
10. Important documents		
11. Understanding booting sequence of beaglebone black hardware		4
12. Preparing SD card for SD boot		
13. Copying boot images to SD card		5
14. Booting BBB via SD card		
15. Making SD boot default on BBB by erasing eMMC MBR		
16. Updating Linux kernel image		7
pre-built image of the kernel : the kernel version is 4.4.62.		
<code>/home/scott/workspace/src/linux_bbb_4.14 3/26 4.14</code>		
<pre> 174 git init 175 git clone https://github.com/beagleboard/linux.git linux_bbb_4.14 178 cd linux_bbb_4.14/ 186 git checkout 4.14.108-ti-r130 443 cd workspace/ 445 cd src 447 cd linux_bbb_4.14/ </pre>		
kernel_image_update_5.10.pdf	5.10-rt	
17. Linux kernel compilation		
<code>/home/scott/workspace/src/linux_bbb_4.14</code>		
18. Modules compilation		

19. Modules install	
20. Update new boot images and modules in SD card	
21. Enabling internet over USB	8

2. Linux Kernel module 1h38m

22. Introduction to Linux kernel module	43~46
23. User space Vs kernel space	89~90
24. LKM writing syntax	47~58

include/linux has all the kernel header files. e.g. linux/modules.h
stdio.h is a user header file.

25. __init and __exit macros	59~62	9
26. LKM entry point registration and other macros	63~66	10
linux_bbb_4.14		
27. Hello World LKM	67~70	11 001
28. Building a Linux kernel module	71~75	12

LKM

static
dynamic
intree

out of tree : kbuild to build modules w/ "a prebuilt kernel source w/ config & headers"
<https://www.kernel.org/doc/Documentation/kbuild/modules.txt>

Two ways to obtain a prebuilt kernel version:

Download kernel from your distributor and build it by yourself
Install the Linux-headers- of the target Linux kernel

make -C \$KDIR -M \$PWD [modules/modules_install/clean/help]

obj-<X> := <module_name>.o -> <module_name>.ko

X = n: do not compile
X = y: compile and link it with kernel
X = m: compile as dynamically loadable kernel module

29. Compilation and testing of an LKM

```
$ uname -r
5.3.0-40-generic

$ make -C /lib/modules
4.14.108/      5.3.0-28-generic/      5.3.0-40-generic/      -> lecture

scott@host:~/workspace/ldd/custom_drivers/001hello_world$ ll /lib/modules/
4.14.108/      5.4.0-150-generic/      5.4.0-84-generic/      -> 18.04 LTS
```

Ubuntu release	Arch	Kernel Version	
-----	-----	-----	
Ubuntu 20.04 LTS	64-bit x86	5.4 (GA)	
Ubuntu 18.04 LTS	64-bit x86	5.4 (HWE)	Hardware Enablement : the most recent versions of the Linux kernel.
Ubuntu 18.04 LTS	64-bit x86	4.15 (GA)	General Availability : the most stable kernel with an original LTS
Ubuntu 16.04 LTS	64-bit x86	4.15 (HWE)	

```
$ make -C /lib/modules/5.3.0-40-generic/build/ M=$PWD modules
$ ls
main.c main.ko main.mod mdin.mod.c main.mod.o main.o Makefile modules.order Module.symvers
$ sudo insmod main.ko
$ dmesg
[164.xxxx] main: loading out-of-tree module tains kernel.
[164.xxxx] main: module verification failed: signature and /or required key missing - tainting kernel
[164.xxxx] Hello world
$ sudo rmmod main.ko
$ dmesg
[288.xxxx] Good bye world
```

30. Testing of an LKM on target

```
$ sudo make ARCH=arm CROSS_COMPILE=arm-linux-gnueabihf- -C /home/scott/workspace/src/linux_4.14/ M=$PWD modules      => 5.3.40 in lecture
$ file main.ko
main.ko: ELF 32-bit LSB relocatable, ARM, EABIS version 1 (SYSV), ....
$ modinfo main.ko
filename:      /path/to/main.ko
board:         Beaglebone black REV A5
description:   a simple hello world kernel module
author:        ME
license:       GPL
depends:
name:          main
vermagic:      4.14.108 SMP preemp mod_unload modversions ARMv7 p2v8
$
$ arm-linux-gnueabihf-objdump -h main.ko
main.ko:      file format elf32-littlearm
```

```

section:
idx name  size      VMA      LMA      file off align
0 .note.
1 .text
2 .init.text
3 .exit.text
4 .ARM.extab.init.text
5 .ARM.exidx.init.text
6 .ARM.extab.exit.text
7 .ARM.exidx.exit.text
8 .modinfo
9 .rodata.str1.4
10 __version
11 .data
12 .gnu.linkonce.this_module
13 .plt
14 .init.plt
15 .bss
16 .comment
17 .note.GNU-stack
18 .ARM.attributes

# sudo insmod main.ko
[ 200.35xxx] main: loading out-of-tree module taints kernel
[ 200.35xxx] Hello world
# dmesg | tail
[ 200.35xxx] main: loading out-of-tree module taints kernel
[ 200.35xxx] Hello world
# sudo rmmod main.ko
[ 288.35xxx] Good bye world

/usr # dmesg      <- my BBB

[ 10.627696] g_ether gadget: Ethernet Gadget, version: Memorial Day 2008
[ 10.639452] g_ether gadget: g_ether ready
[ 10.923157] g_ether gadget: high-speed config #2: RNDIS
[ 121.502878] random: crng init done
[ 215.820416] main: loading out-of-tree module taints kernel.
[ 215.826708] Hello world
[ 226.387175] Good bye world
/usr #

```

<https://www.kernel.org/doc/Documentation/kbuild/kconfig-language.txt>)

You have to add the Linux kernel module inside the Linux kernel source tree and let the Linux build system builds that. If you want to list your kernel module selection in kernel menuconfig, then create and use a Kconfig file

. Kconfig file

```
scott@vbox:~/.../src/linux_bbb_4.14/drivers/char/my_c_dev $ ls
Kconfig      main.c
```

<= newly added

```
    menu "my custom modules"
        config CUSTOM_HELLOWORLD
            tristate "hello world module support"
            default m
    endmenu
```

```
scott@vbox:~/.../src/linux_bbb_4.14/drivers/char $ vi Kconfig
```

<= upper level Kconfig

```
...
source "drivers/char/my_c_dev/Kconfig"

endmenu
```

```
scott@vbox:~/.../src/linux_bbb_4.14/drivers/char/my_c_dev $ vi Makefile
```

```
#obj-<config_item> += <module>.o

obj-$(CONFIG_CUSTOM_HELLOWORLD) += main.o
it's y or n or m.
```

<= newly added
// \$(CONFIG_CUSTOM_HELLOWORLD) if we don't know

```
scott@vbox:~/.../src/linux_bbb_4.14/drivers/char $ vi Makefile
```

<= upper level Makefile

```
...
obj-y
```

+= my_c_dev/

// To select a kernel module under that menu, we have config item. But, select our folder we don't have any config item.

```
scott@vbox:~/.../src/linux_bbb_4.14 $ make ARCH=arm menuconfig
```

```
Device Drivers --->
```

```
Character devices --->
```

```
fastbit custom modules --->
```

```
<m> helloworld module support : dynamically linked. '*' means statically linked.
```

```
<help>
```

```
There is no help available for this option.
```

```
Symbol: CUSTOM_HELLOWORLD [=m]
```

```
Type : tristate
```

```
Prompt: helloworld module support
```

```
Location:
```

```
-> Device Drivers
```

```
-> Character devices
```

```
-> fastbit custom modules
```

```
Defined at drivers/char/my_c_dev/Kconfig:2
```

```
scott@vbox:~/.../src/linux_bbb_4.14 $ make ARCH=arm CROSS_COMPILE=arm-linux-gnueabihf- modules -j4
```

```
scripts/kconfig/conf --silentoldconfig Kconfig
```

```
...
```

```
CC [M] drivers/char/my_c_dev/main.o
```

```
Building modules, stage 2.
```

```
..
```

```
LD [M] drivers/char/my_c_dev/main.ko
```

```
scott@vbox:~/.../src/linux_bbb_4.14/drivers/char/my_c_dev $ modinfo main.ko
```

```
filename: /home/scott/workspace/src/linux_bbb_4.14/drivers/char/my_c_dev/main.ko
```

```
board: Beaglebone black REV A5
```

```
description: A simple hello world kernel module
```

```
author: Kiran Nayak
```

```
license: GPL
```

```
depends:
```

```
intree: Y <= !!!!!
```

```
name: main
```

```
vermagic: 4.14.108+ SMP preempt mod_unload modversions ARMv7 p2v8
```

33. printk : to revisit

3. Character device and driver

2h5m

34. What is device driver ?

88

```
write(fd, 0xab);    // echo 0xAB > /dev/rtc
```

Udev : populates /dev w/ device files

devices files are under VFS and identified by major/minor numbers

35. A char driver, char device and char device number

15

Let's create a device number.

Your driver has to ask the kernel to dynamically allocate the device number or numbers.

Basically, what you should be doing here is, you should be using kernel APIs and kernel utilities (shown in Figure 6) in order to request various services from the kernel.

Now, to create a device number, you just have to use a kernel API `alloc_chrdev_region()`.

So, you have to use `alloc_chrdev_region()`. This creates a device number. And for the registration, you can use these APIs `cdev_init()` and `cdev_add()`.

And after that, the driver should create a device files. For that, you can use these kernel APIs `class_create()` and `device_create()`.

The creation things we are going to do in module initialization function.

Whenever you load a module, so these creation services must be executed and your driver must be ready to accept a system calls from the user space program.

That's why it makes sense to do this creation process in the module initialization function.

When you remove the module, it's better you delete all those resources what you requested from the kernel.

Otherwise, it will simply consume a resources of the kernel. That's why, let's say, if you use `unregister_chrdev_region()`, it will delete the device number which is allocated for your module, that it can be reused for some other module.

After that, you can use `cdev_del()` to delete the registration, that will free some memory. `class_destroy` and `device_destroy` will delete your device files.

That's why, all these deletion things we are going to do in module clean-up function. Because these deletion things should be executed whenever you remove the module.

36. Dynamically allocating char device numbers	108	16
* 37. Pseudo character driver implementation		17
38. Character device registration	114	

```

scott@host:~/workspace/ldd/custom_drivers/custom_drivers/001hello_world$ nm main.ko
0000000000000000 T cleanup_module
                U __fentry__
0000000000000000 t helloworld_cleanup
0000000000000000 t helloworld_init
0000000000000000 T init_module
0000000000000000 r _note_6
                U printk
0000000000000000 D __this_module          <= !!!
000000000000004d r __UNIQUE_ID_author37
0000000000000000 r __UNIQUE_ID_board39
...
00000000000000ae r __UNIQUE_ID_vermagic36
0000000000000000 r ____versions

```

```
>> main.mod.c
```

```

#include <linux/build-salt.h>
#include <linux/module.h>
#include <linux/vermagic.h>
#include <linux/compiler.h>

```

```
BUILD_SALT;
```

```

MODULE_INFO(vermagic, VERMAGIC_STRING);
MODULE_INFO(name, KBUILD_MODNAME);

```

```

__visible struct module __this_module          <= !!!
__section(.gnu.linkonce.this_module) = {
    .name = KBUILD_MODNAME,

```

```

struct file_operations {
    struct module *owner;                <- __this_module : to prevent a module from being unloaded while the
    structure in use
    loff_t (*llseek) (struct file *, loff_t, int);
}

```


cdev_init() -> just initialization

39. Character device registration contd. 121

cdev_add() -> real registration

40. Character driver file operation methods 18

Character Driver System Calls

What is a file object?

Whenever a file is opened, a file object is created in the kernel space.

The file object stores information about the interaction between an open file and a user process.

def_chr_fops

open		user

-> do_sys_open	--> return 'fd' to user space (-1 means failure)	kernel
-> do_filp_open	--> 'file' object allocation	
-> do_dentry_open	--> using default (dummy) fops, e.g. chrdev_open for open	
-> chrdev_open	--> replacing default fops	
-> pcd_open		

do_filp_open : a file object is created

ref) understanding linux kernel 3ed, p524

When device file gets created

- 1) create device file using udev
- 2) inode object gets created in memory and inode's i_rdev field is initialized with device number(dev_t)
- 3) inode object's i_fop field is set to dummy default file operations (def_chr_fops)

When user process executes open system call

- 1) user invokes open system call on the device file
- 2) file object gets created (do_filp_open)
- 3) inode's i_fop gets copied to file object's f_op (dummy default file operations of char device file)
- 4) open function of dummy default file operations gets called (chrdev_open)
- 5) inode object's i_cdev field is initialized with cdev which you added during cdev_add (lookup happens using inode->i_rdev field)

- 6) inode->cdev->fops (this is a real file operations of your driver) gets copied to file->f_op
7) file->f_op->open method gets called (read open method of your driver)

```
static int chrdev_open(struct inode *inode, struct file *filp)
{
    const struct file_operations *fops;
    struct cdev *p;
    struct cdev *new = NULL;
    int ret = 0;
    spin_lock(&cdev_lock);
    p = inode->i_cdev;

    if (!p) {
        struct kobject *kobj;
        int idx;
        spin_unlock(&cdev_lock);
        kobj = kobj_lookup(cdev_map, inode->i_rdev, &idx);
        if (!kobj) return -ENXIO;
        new = container_of(kobj, struct cdev, kobj);
        spin_lock(&cdev_lock);
        p = inode->i_cdev;
        if (!p) {
            inode->i_cdev = p = new;
            list_add(&inode->i_devices, &p->list);
            new = NULL;
        } else if (!cdev_get(p))
            ret = -ENXIO;
    } else if (!cdev_get(p))
        ret = -ENXIO;

    spin_unlock(&cdev_lock);
    cdev_put(new);

    if (ret) return ret;
    ret = -ENXIO;
    fops = fops_get(p->ops);

    if (!fops) goto out_cdev_put;
    replace_fops(filp, fops);

    if (filp->f_op->open) {
        ret = filp->f_op->open(inode, filp);
        if (ret) goto out_cdev_put;
    }
    return 0;
}

out_cdev_put:
cdev_put(p);
```

```

        return ret;
    }

```

41. Character driver file operation methods contd.

Open Method	Close system call
Release Method:	
Write Method:	
Llseek Method:	

42. Implementing file operation methods

Accessing File Operations Structure

```

include/linux/fs.h

struct file_operations {
    struct module *owner;
    loff_t (*llseek) (struct file *, loff_t, int);
    ssize_t (*read) (struct file *, char __user *, size_t, loff_t *);
    ssize_t (*write) (struct file *, const char __user *, size_t, loff_t *);
    ssize_t (*read_iter) (struct kiocb *, struct iov_iter *);
    ...
    int (*clone_file_range)(struct file *, loff_t, struct file *, loff_t, u64);
    ssize_t (*dedupe_file_range)(struct file *, u64, u64, struct file *, u64);
} __randomize_layout;

```

Implementing File Operation Functions

Method Name Replacement
 Creating Method Variables
 Adding Return Values

43. File operations structure initialization

```

CarBMW = {a, b, c, d}; // C89
CarBMW = {.A = a, .B = b, .C = c, .D = d}; // C99

```

```
class_create and
device_create
```

45. Character driver cleanup function implementation

```
scott@host:~/workspace/.../002pseudo_char_driver$ sudo insmod pcd.ko
```

```
[46707.310503] pcd_driver_init :Device number <major>:<minor> : 240:0
[46707.310795] pcd_driver_init :Module init was successful
```

```
scott@host:~/workspace/.../002pseudo_char_driver$
```

```
scott@host:~/workspace/.../002pseudo_char_driver$ cd /sys/class/pc
```

```
pcd_class/ pci_bus/ pci_epc/
```

```
scott@host:~/workspace/.../002pseudo_char_driver$ ls /sys/class/pcd_class
```

```
pcd
```

```
scott@host:~/workspace/.../002pseudo_char_driver$ ls /sys/class/pcd_class/pcd
```

```
dev power subsystem uevent
```

```
scott@host:/sys/class/pcd_class/pcd$ cat dev
```

```
240:0
```

```
scott@host:/sys/class/pcd_class/pcd$ cat uevent
```

```
MAJOR=240
MINOR=0
DEVNAME=pcd
```

```
scott@host:/sys/class/pcd_class/pcd$ ll /dev/pcd
```

```
crw----- 1 root root 240, 0 Apr 24 19:36 /dev/pcd
```

```
scott@host:/sys/class$ ls
```

```

ata_device  devcoredump  firmware  iommu      ^^^^^^^^^^
ata_link    devfreq      gpio      leds       pcd_class
ata_port    devfreq-event graphics  mdio_bus  pci_bus
backlight   dma          hidraw    mem        pci_epc
bdi         dmi          hwmon     misc       phy
            dmidec      i2c       pinctrl    pps
            dmidec      i2c       pinctrl    printer
            dmidec      i2c       pinctrl    ptp
            dmidec      i2c       pinctrl    pwm
            dmidec      i2c       pinctrl    rapidio_port
            dmidec      i2c       pinctrl    rtc
            dmidec      i2c       pinctrl    scsi_device
            dmidec      i2c       pinctrl    scsi_disk
            dmidec      i2c       pinctrl    scsi_generic
            dmidec      i2c       pinctrl    scsi_host
            dmidec      i2c       pinctrl    thermal
            dmidec      i2c       pinctrl    tpm
            dmidec      i2c       pinctrl    tpmrm
            dmidec      i2c       pinctrl    tty
            dmidec      i2c       pinctrl    vc
            dmidec      i2c       pinctrl    watchdog
            dmidec      i2c       pinctrl    wakeup
```

block	drm	i2c-adapter	mmc_host	power_supply	regulator	sound	vfio
bsg	drm_dp_aux_dev	i2c-dev	nd	ppdev	remoteproc	spi_master	virtio-ports
dax	extcon	input	net	ppp	rkill	spi_slave	vtconsole

```
scott@host:/sys/class$ sudo rmmod pcd.ko
scott@host:/sys/class$ dmesg
```

```
[46707.310503] pcd_driver_init :Device number <major>:<minor> : 240:0
[46707.310795] pcd_driver_init :Module init was successful
[47219.531427] pcd_driver_cleanup :module unloaded
```

4. Character driver file operation implementation 1h1m 134~

46. Understanding read method
 47. Understanding error codes
 48. Read method implementation

49. Understanding write method
 50. Write method implementation

51. lseek method
 52. lseek method implementation

53. Testing pseudo char driver

```
$ sudo -s
# insmod pcd.ko
# dmesg
```

```
[54292.379629] pcd_driver_init :Device number <major>:<minor> : 240:0
[54292.379692] pcd_driver_init :Module init was successful
```

```
# echo "Hello, welcome to the course" > /dev/pcd
```

```
[54292.379629] pcd_driver_init :Device number <major>:<minor> : 240:0
[54292.379692] pcd_driver_init :Module init was successful
```

```
[54432.442779] pcd_open :open was successful -> echo issued an OPEN system cal.
```

```
[54432.442787] pcd_write :write requested for 29 bytes -> echo issued a WRITE system cal.
[54432.442788] pcd_write :current file position = 0
[54432.442788] pcd_write :Number of bytes successfully write = 29
[54432.442789] pcd_write :Updated file position = 29
```

```

[54432.442791] pcd_release :release was successful                                -> echo issued an RELEASE system cal.

# cat /dev/pcd
Hello, welcome to the course

[54679.429313] pcd_open :open was successful

[54679.429319] pcd_read :read requested for 131072 bytes
[54679.429319] pcd_read :current file position = 0
[54679.429321] pcd_read :Number of bytes successfully read = 512
[54679.429321] pcd_read :Updated file position = 512

[54679.429324] pcd_read :read requested for 131072 bytes
[54679.429325] pcd_read :current file position = 512
[54679.429325] pcd_read :Number of bytes successfully read = 0                -> end of file
[54679.429326] pcd_read :Updated file position = 512

[54679.429331] pcd_release :release was successful

# cp test.txt /dev/pcd

root@host:~/workspace/ldd/custom_drivers/custom_drivers/002pseudo_char_driver# cp pcd.c /dev/pcd
cp: error writing '/dev/pcd': No space left on device

root@host:~/workspace/ldd/custom_drivers/custom_drivers/002pseudo_char_driver# dmesg | tail - 20

[55108.260918] pcd_open :open was successful

[55108.260928] pcd_write :write requested for 4289 bytes
[55108.260929] pcd_write :current file position = 0
[55108.260929] pcd_write :Number of bytes successfully write = 512
[55108.260930] pcd_write :Updated file position = 512

[55108.260931] pcd_write :write requested for 3777 bytes                        4289 - 512 = 3777
[55108.260931] pcd_write :current file position = 512
[55108.260932] pcd_write :Number of bytes successfully write = 0                -> end of memory
[55108.260932] pcd_write :Updated file position = 512

[55108.261188] pcd_release :release was successful

root@host:~/workspace/ldd/custom_drivers/custom_drivers/002pseudo_char_driver# cp pcd.c /dev/pcd
cp: error writing '/dev/pcd': Cannot allocate memory

root@host:~/workspace/ldd/custom_drivers/custom_drivers/002pseudo_char_driver# dmesg | tail -10

[55328.641690] pcd_driver_init :Module init was successful

```

```

[55338.920580] pcd_open :open was successful

[55338.920590] pcd_write :write requested for 4372 bytes
[55338.920590] pcd_write :current file position = 0
[55338.920591] pcd_write :Number of bytes successfully write = 512
[55338.920592] pcd_write :Updated file position = 512

[55338.920593] pcd_write :write requested for 3860 bytes
[55338.920593] pcd_write :current file position = 512
[55338.920593] pcd_write :--No space left on the device          <-    count == 0

[55338.920812] pcd_release :release was successful

```

54. Error handling

5. Char driver with multiple device nodes 1h49m 187~

55. pcd driver with multiple devices

- 56. Pcd driver with multiple devices code implementation part-1
- 57. Pcd driver with multiple devices code implementation part-2
- 58. Pcd driver with multiple devices code implementation part-3
- 59. Pcd driver with multiple devices code implementation part-4
- 60. Pcd driver with multiple devices code implementation part-5
- 61. Pcd driver with multiple devices code implementation part-6
- 62. Pcd driver with multiple devices code implementation part-7

63. Pcd driver with multiple devices testing

```
$ sudo insmod pcd_n.ko
```

```

[ 3708.848439] pcd_driver_init : Device number <major>:<minor> = 240:0
[ 3708.848484] pcd_driver_init : Device number <major>:<minor> = 240:1
[ 3708.848499] pcd_driver_init : Device number <major>:<minor> = 240:2
[ 3708.848514] pcd_driver_init : Device number <major>:<minor> = 240:3
[ 3708.848535] pcd_driver_init : Module init was successful

```

```
$ echo hello > /dev/pcdev-1
```

```
bash: /dev/pcdev-1: Permission denied
```

```
[ 3708.848535] pcd_driver_init : Module init was successful  
[ 3708.8485XX] pcd_open : minor access = 0  
[ 3708.8485XX] pcd_open : open was unsuccessful
```

```
$ strace dd if=pcd_n.c of=/dev/pcdev-1
```

```
execve("/bin/dd", ["dd", "if=pcd_n.c", "of=/dev/pcdev-1"], 0x7ffd37eb4840 /* 53 vars */) = 0  
brk(NULL) = 0x55e06a28a000  
access("/etc/ld.so.nohwcap", F_OK) = -1 ENOENT (No such file or directory)  
access("/etc/ld.so.preload", R_OK) = -1 ENOENT (No such file or directory)  
openat(AT_FDCWD, "/etc/ld.so.cache", O_RDONLY|O_CLOEXEC) = 3  
fstat(3, {st_mode=S_IFREG|0644, st_size=79287, ...}) = 0  
mmap(NULL, 79287, PROT_READ, MAP_PRIVATE, 3, 0) = 0x7f18eb28a000  
close(3) = 0  
access("/etc/ld.so.nohwcap", F_OK) = -1 ENOENT (No such file or directory)  
openat(AT_FDCWD, "/lib/x86_64-linux-gnu/libc.so.6", O_RDONLY|O_CLOEXEC) = 3  
read(3, "\177ELF\2\1\1\3\0\0\0\0\0\0\0\3\0\>\0\1\0\0\0\240\35\2\0\0\0\0\0"... , 832) = 832  
fstat(3, {st_mode=S_IFREG|0755, st_size=2030928, ...}) = 0  
mmap(NULL, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0x7f18eb288000  
mmap(NULL, 4131552, PROT_READ|PROT_EXEC, MAP_PRIVATE|MAP_DENYWRITE, 3, 0) = 0x7f18eac84000  
mprotect(0x7f18eae6b000, 2097152, PROT_NONE) = 0  
mmap(0x7f18eb06b000, 24576, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|MAP_DENYWRITE, 3, 0x1e7000) = 0x7f18eb06b000  
mmap(0x7f18eb071000, 15072, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|MAP_ANONYMOUS, -1, 0) = 0x7f18eb071000  
close(3) = 0  
arch_prctl(ARCH_SET_FS, 0x7f18eb289540) = 0  
mprotect(0x7f18eb06b000, 16384, PROT_READ) = 0  
mprotect(0x55e0692dd000, 4096, PROT_READ) = 0  
mprotect(0x7f18eb29e000, 4096, PROT_READ) = 0  
munmap(0x7f18eb28a000, 79287) = 0  
rt_sigaction(SIGINT, NULL, {sa_handler=SIG_DFL, sa_mask=[], sa_flags=0}, 8) = 0  
rt_sigaction(SIGUSR1, {sa_handler=0x55e0690d00e0, sa_mask=[INT_USR1], sa_flags=SA_RESTORER, sa_restorer=0x7f18eacc2f10}, NULL, 8) = 0  
rt_sigaction(SIGINT, {sa_handler=0x55e0690d00d0, sa_mask=[INT_USR1], sa_flags=SA_RESTORER|SA_NODEFER|SA_RESETHAND, sa_restorer=0x7f18eacc2f10}, NULL, 8) = 0  
brk(NULL) = 0x55e06a28a000  
brk(0x55e06a2ab000) = 0x55e06a2ab000  
openat(AT_FDCWD, "/usr/lib/locale/locale-archive", O_RDONLY|O_CLOEXEC) = 3  
fstat(3, {st_mode=S_IFREG|0644, st_size=3004224, ...}) = 0  
mmap(NULL, 3004224, PROT_READ, MAP_PRIVATE, 3, 0) = 0x7f18ea9a6000  
close(3) = 0  
openat(AT_FDCWD, "pcd_n.c", O_RDONLY) = 3  
dup2(3, 0) = 0  
close(3) = 0  
lseek(0, 0, SEEK_CUR) = 0  
openat(AT_FDCWD, "/dev/pcdev-1", O_WRONLY|O_CREAT|O_TRUNC, 0666) = -1 EACCES (Permission denied)  
openat(AT_FDCWD, "/usr/share/locale/locale.alias", O_RDONLY|O_CLOEXEC) = 3  
fstat(3, {st_mode=S_IFREG|0644, st_size=2995, ...}) = 0
```



```

read(3, "# Locale name alias data base.\n#"... , 4096) = 2995
read(3, "", 4096) = 0
close(3) = 0
openat(AT_FDCWD, "/usr/share/locale/en_US.UTF-8/LC_MESSAGES/coreutils.mo", O_RDONLY) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/usr/share/locale/en_US.utf8/LC_MESSAGES/coreutils.mo", O_RDONLY) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/usr/share/locale/en_US/LC_MESSAGES/coreutils.mo", O_RDONLY) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/usr/share/locale/en.UTF-8/LC_MESSAGES/coreutils.mo", O_RDONLY) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/usr/share/locale/en.utf8/LC_MESSAGES/coreutils.mo", O_RDONLY) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/usr/share/locale/en/LC_MESSAGES/coreutils.mo", O_RDONLY) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/usr/share/locale-langpack/en_US.UTF-8/LC_MESSAGES/coreutils.mo", O_RDONLY) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/usr/share/locale-langpack/en_US.utf8/LC_MESSAGES/coreutils.mo", O_RDONLY) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/usr/share/locale-langpack/en_US/LC_MESSAGES/coreutils.mo", O_RDONLY) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/usr/share/locale-langpack/en.UTF-8/LC_MESSAGES/coreutils.mo", O_RDONLY) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/usr/share/locale-langpack/en.utf8/LC_MESSAGES/coreutils.mo", O_RDONLY) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/usr/share/locale-langpack/en/LC_MESSAGES/coreutils.mo", O_RDONLY) = 3
fstat(3, {st_mode=S_IFREG|0644, st_size=578, ...}) = 0
mmap(NULL, 578, PROT_READ, MAP_PRIVATE, 3, 0) = 0x7f18eb29d000
close(3) = 0
write(2, "dd: ", 4dd: ) = 4
write(2, "failed to open '/dev/pcdev-1'", 29failed to open '/dev/pcdev-1') = 29
openat(AT_FDCWD, "/usr/share/locale/en_US.UTF-8/LC_MESSAGES/libc.mo", O_RDONLY) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/usr/share/locale/en_US.utf8/LC_MESSAGES/libc.mo", O_RDONLY) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/usr/share/locale/en_US/LC_MESSAGES/libc.mo", O_RDONLY) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/usr/share/locale/en.UTF-8/LC_MESSAGES/libc.mo", O_RDONLY) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/usr/share/locale/en.utf8/LC_MESSAGES/libc.mo", O_RDONLY) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/usr/share/locale/en/LC_MESSAGES/libc.mo", O_RDONLY) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/usr/share/locale-langpack/en_US.UTF-8/LC_MESSAGES/libc.mo", O_RDONLY) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/usr/share/locale-langpack/en_US.utf8/LC_MESSAGES/libc.mo", O_RDONLY) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/usr/share/locale-langpack/en_US/LC_MESSAGES/libc.mo", O_RDONLY) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/usr/share/locale-langpack/en.UTF-8/LC_MESSAGES/libc.mo", O_RDONLY) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/usr/share/locale-langpack/en.utf8/LC_MESSAGES/libc.mo", O_RDONLY) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/usr/share/locale-langpack/en/LC_MESSAGES/libc.mo", O_RDONLY) = -1 ENOENT (No such file or directory)
write(2, ": Permission denied", 19: Permission denied) = 19
write(2, "\n", 1
) = 1
close(2) = 0
exit_group(1) = ?
+++ exited with 1 +++

```

```
scott@host:~/workspace/ldd/ex/003_pseudo_char_driver_multiple$ sudo dd if=pcd_n.c of=/dev/pcdev-2
```

```

dd: writing to '/dev/pcdev-2': Cannot allocate memory
2+0 records in
1+0 records out
512 bytes copied, 0.000306726 s, 1.7 MB/s

```

```
scott@host:~/workspace/ldd/ex/003_pseudo_char_driver_multiple$ dmesg | tail
```

```
[ 4143.773984] pcd_open : minor access = 1
[ 4143.773985] pcd_open : open was successful

[ 4143.773999] pcd_write : Write requested for 512 bytes
[ 4143.774000] pcd_write : Current file position = 0
[ 4143.774001] pcd_write : Number of bytes successfully written = 512
[ 4143.774002] pcd_write : Updated file position = 512

[ 4143.774004] pcd_write : Write requested for 512 bytes
[ 4143.774005] pcd_write : Current file position = 512
[ 4143.774006] pcd_write : No space left on the device

[ 4143.774289] pcd_release : release was successful
```

```
$ sudo dd if=pcd_n.c of=/dev/pcdev-2 count=1
```

```
1+0 records in
1+0 records out
512 bytes copied, 0.00010661 s, 4.8 MB/s
```

```
$ dmesg | tail -15
```

```
[ 4336.492774] pcd_open : minor access = 1
[ 4336.492775] pcd_open : open was successful

[ 4336.492784] pcd_write : Write requested for 512 bytes
[ 4336.492785] pcd_write : Current file position = 0
[ 4336.492785] pcd_write : Number of bytes successfully written = 512
[ 4336.492786] pcd_write : Updated file position = 512

[ 4336.492788] pcd_release : release was successful
```

```
scott@host:~/workspace/ldd/ex/003_pseudo_char_driver_multiple$ sudo dd if=pcd_n.c of=/dev/pcdev-3 count=1 bs=100
```

```
1+0 records in
1+0 records out
100 bytes copied, 0.000108374 s, 923 kB/s
```

```
scott@host:~/workspace/ldd/ex/003_pseudo_char_driver_multiple$ sudo cat /dev/pcdev-3
```

```
#include<linux/module.h>
#include<linux/fs.h>
#include<linux/cdev.h>
#include<linux/device.h>
#include<scott@host:~/workspace/ldd/ex/003_pseudo_char_driver_multiple$
```

```
scott@host:~/workspace/ldd/ex/003_pseudo_char_driver_multiple$ dmesg | tail -15
```

```
[ 4628.650774] pcd_open : minor access = 2
[ 4628.650775] pcd_open : open was successful

[ 4628.650781] pcd_read : Read requested for 131072 bytes
[ 4628.650782] pcd_read : Current file position = 0
[ 4628.650783] pcd_read : Number of bytes successfully read = 1024
[ 4628.650784] pcd_read : Updated file position = 1024

[ 4628.650794] pcd_read : Read requested for 131072 bytes
[ 4628.650794] pcd_read : Current file position = 1024
[ 4628.650795] pcd_read : Number of bytes successfully read = 0
[ 4628.650795] pcd_read : Updated file position = 1024

[ 4628.650801] pcd_release : release was successful
```

64. Pcd driver with multiple devices testing contd

65. Pcd driver with multiple devices lseek implementation

66. Container of discussion

6. Platform bus, devices and drivers 2h39m 209

67. Platform devices and drivers 209 ~ 225

```
linux-3.16.84/arch/arm/mach-omap2/board-xxxxxx.h
platform_add_devices(devkit8000_devices, ARRAY_SIZE(devkit8000_devices));
```

```
linux-4.14/arch/arm/mach-omap2/board-generic.h
```

68. Example of platform drivers 226 ~ 228

p227 - all platform devices

p228 - platform drivers = controller drivers = bus drivers

a device and a controller : A controller always controls a device

The controller drivers already available given by the SOC vendor.

```
#define platform_driver_register(drv) __platform_driver_register(drv, THIS_MODULE)
extern int __platform_driver_register(struct platform_driver *, struct module *);
```

Platform driver structure : struct platform_driver

Registering a platform device : int platform_device_register(struct platform_device *pdev);

Platform device structure : struct platform_device

probe(),	Called when matched platform device is found
remove(),	
shutdown(),	Called at shut-down time to quiesce the device
suspend(),	Called to put the device to sleep mode. Usually to a low power state
resume(), ...	Called to bring a device from sleep mode

Platform device - driver matching 234

- "matching" mechanism of the bus core

The Linux platform core implementation maintains platform device and driver lists. Whenever you add a new platform device or driver, this list gets updated and matching mechanism triggers.

Every bus type has its match function, where the device and driver list will be scanned.

Points to remember 239

- Whenever a new device or a new driver is added, the matching function of the platform bus runs, and if it finds a matching platform device for a platform driver, the probe function of the matched driver will get called. Inside the probe function, the driver configures the detected device.
- Details of the matched platform device will be passed to the probe function of the matched driver so that driver can extract the platform data and configure it.

Probe function of the platform driver 241

- Probe function must be implemented by the platform driver and should be registered during platform_driver_register().
- When the bus matching function detects the matching device and driver, probe function of the driver gets called with detected platform device as an input argument
- Note that probe() should in general, verify that the specified device hardware actually exists. Sometimes platform setup code can't be sure. The probing can use device resources, including clocks, and device platform_data.
- The the probe function is responsible for
- Device detection and initialization
- Allocation of memories for various data structures,
- Mapping i/o memory
- Registering interrupt handlers

- Registering device to kernel framework, user level access point creations, etc
- The probe may return 0(Success) or error code. If probe function returns a non-zero value, meaning probing of a device has failed.

Remove function of the platform driver 242

- Remove function gets called when a platform device is removed from the kernel to unbind a device from the driver or when the kernel no longer uses the platform device
- Remove function is responsible for
- Unregistering the device from the kernel framework
- Free memory if allocated on behalf of a device
- Shutdown/De-initialize the device

<https://www.kernel.org/doc/Documentation/driver-model/platform.txt>

70. Platform driver code exercise 245

essential functions

```
int pcd_platform_driver_remove(struct platform_device *pdev)
int pcd_platform_driver_probe(struct platform_device *pdev)

static int __init pcd_platform_driver_init(void)
{
    platform_driver_register();
    return 0;
}

static void __exit pcd_platform_driver_exit(void)
{
    platform_driver_unregister();
}
```

Code exercise: 248

- Implementation of pseudo character driver as platform driver
- Repeat the exercise pseudo character driver with multiple devices as a platform driver.
 - The driver should support multiple pseudo character devices(pcdevs) as platform devices
 - Create device files to represent platform devices
 - The driver must give open, release, read, write, lseek methods to deal with the devices

kernel module 1 : platform driver 250
kernel module 2 : platform device setup

1. Create 2 platform devices and initialize them with required information
 - Name of a platform device
 - Platform data
 - Id of the device
 - Release function for the device
2. Register platform devices with the Linux kernel

71. Platform device setup code implementation

```

struct pcdev_private_data      /*Device private data structure */
{
    struct pcdev_platform_data pdata;
    char *buffer;
    dev_t dev_num;
    struct cdev cdev;
};

struct pcdrv_private_data      /*Driver private data structure */
{
    int total_devices;
    dev_t device_num_base;
    struct class *class_pcd;
    struct device *device_pcd;
};

struct platform_device          347
{
    const char *name;           <= used to match
    int id;                     <= for multiple instances
    struct device dev;          311
}

```

72. Platform device setup code implementation contd. : setup device(module 2)

Platform device release function 251

- Callback to free the device after all references have gone away.
- This should be set by the allocator of the device

```

platform_add_devices(platform_pcdevs,ARRAY_SIZE(platform_pcdevs) );
platform_device_unregister(&platform_pcdev_1);

```

```
$ sudo insmod pcd_device_setup.ko
```

```
$ ls /sys/devices/platform/
eisa.0          pcdev-A1x.0      platform-framebuffer.0  uevent
'Fixed MDIO bus.0'  pcdev-B1x.1      power
i8042          pcdev-C1x.2      reg-dummy
intel_rapl_msr.0  pcdev-D1x.3      rtc_cmos
kgdboc         pcspkr           serial8250
```

```
$ dmesg | tail -30
[18617.022770] pcdev_platform_init : Device setup module loaded
```

```
$ sudo rmmod pcd_device_setup.ko
$ dmesg | tail -30
[18861.674563] pcdev_release : Device released
[18861.674825] pcdev_release : Device released
[18861.674869] pcdev_release : Device released
[18861.674880] pcdev_release : Device released
[18861.674881] pcdev_platform_exit : Device setup module unloaded
```

```
$ ls /sys/devices/platform/
eisa.0          pcspkr           serial8250
'Fixed MDIO bus.0'  platform-framebuffer.0  ueven
i8042          power
intel_rapl_msr.0  reg-dummy
kgdboc         rtc_cmos
```

73. Platform driver code implementation part-1 : driver(module 1)

```
struct device_driver      in linux/device.h
platform_driver_register  in platform_device.h
```

74. Platform driver code implementation part-2

```
# insmod pcd_device_setup.ko
[18861.674881] pcdev_platform_init : Device setup module loaded

# insmod pcd_platform_driver.ko
[18861.674881] pcdev_platform_init : Device setup module loaded
[18861.674881] pcd_platform_driver_probe : A device is detected
[18861.674881] pcd_platform_driver_init : pcd platform driver loaded

# rmmod pcd_platform_driver.ko
[18861.674881] pcdev_platform_init : Device setup module loaded
[18861.674881] pcd_platform_driver_probe : A device is detected
[18861.674881] pcd_platform_driver_init : pcd platform driver loaded
[18861.674881] pcd_platform_driver_remove : A device is removed
[18861.674881] pcd_platform_driver_cleanup : pcd platform driver unloaded
```

```
# rmmod pcd_device_setup
[18861.674881] pcdev_platform_init : Device setup module loaded
[18861.674881] pcd_platform_driver_probe : A device is detected
[18861.674881] pcd_platform_driver_init : pcd platform driver loaded
[18861.674881] pcd_platform_driver_remove : A device is removed
[18861.674881] pcd_platform_driver_cleanup : pcd platform driver unloaded

[18861.674881] pcdev_release : Device released
[18861.674881] pcdev_release : Device released
[18861.674881] pcdev_platform_exit : Device setup module unloaded
```

or

```
# insmod pcd_platform_driver.ko
[18861.674881] pcd_platform_driver_init : pcd platform driver loaded

# insmod pcd_device_setup.ko
[18861.674881] pcd_platform_driver_init : pcd platform driver loaded
[18861.674881] pcd_platform_driver_probe : A device is detected
[18861.674881] pcdev_platform_init : Device setup module loaded

# rmmod pcd_device_setup
[18861.674881] pcd_platform_driver_init : pcd platform driver loaded
[18861.674881] pcd_platform_driver_probe : A device is detected
[18861.674881] pcdev_platform_init : Device setup module loaded
[18861.674881] pcd_platform_driver_remove : A device is removed
[18861.674881] pcdev_release : Device released
[18861.674881] pcdev_release : Device released
[18861.674881] pcdev_platform_exit : Device setup module unloaded
```

75. Platform driver code implementation part-3

76. Platform driver code implementation part-4

```
pdata = pdev->dev.platform_data;
or pdata = (struct pcdev_platform_data*)dev_get_platdata(&pdev->dev);
```

Kernel memory allocation APIs 325

- kmalloc ()
- kfree ()

```
void* kmalloc( size_t size, gfp_t flags);//include/linux/slab.h 326
```

Used to allocate memory in kernel space by drivers and kernel functions
Memory obtained are physically(RAM) contiguous

gfp_t : get free pages

kzalloc - allocate memory. The memory is set to zero.

77. Platform driver code implementation part-5

78. Platform driver code implementation part-6

~ 335

to free allocated memory by kalloc()

```
struct platform_device {
    ...
    struct device    dev;
    ...
}

struct device {
    ...
    void *platform_data;    <=
    void *driver_data;      <=
    ...
}

// save the device private data pointer in platform device structure
pdev->dev.driver_data = dev_data;          or
dev_set_drvdata(&pdev->dev, dev_data);

==>

struct pcddev_private_data *dev_data = dev_get_drvdata(&pdev->dev);
kfree(dev_data->buffer);
kfree(dev_data);
```

79. Testing platform driver

```
# insmod pcd_device_setup.ko
# insmod pcd_platform_driver.ko
# lsmod                                -> list all the modules
Module                Size      Used by
pcd_platform_driver    16384     0      <= !!
pcd_device_setup       16384     0
...

# rmmod pcd_device_setup
#
# rmmod pcd_platform_driver
```

```
# insmod pcd_platform_driver.ko
# insmod pcd_device_setup.ko
#
# rmmod pcd_platform_driver
```

80. Linux device resource managed functions

336~340

```
devm_kzalloc()      <=>    kalloc & kfree      <= based on the existence of struct device
devm_gpio_get()     <=>    gpio_get()/put()
devm_request_irq()  <=>    request_irq()/free_irq()
```

www.kernel.org/doc/Documentation/driver-model/devres.txt

MFD, MUX, PCI, PHY, MEM, PWM, MDIO, IOMAP, INPUT, IRQ, IO region, IIC, CLOCK, DRM, GPIO, ...

* 81. Using device resource managed kernel functions

```
kxxxx() -> devm_kxxxx()
```

* 82. Testing with more platform devices

```
int platform_add_devices(struct platform_device **devs, int num) {
    for (int i = 0; i < num; i++) {
        platform_device_register(devs[i]);
    }
}
```

kernel crash w/ same ids

```
struct platform_device platform_pcdev_1 = {
    .name = "pcdev-A1x",
    .id = 0,
};
```

```
struct platform_device platform_pcdev_2 = {
    .name = "pcdev-A1x",
    .id = 0,
};
```

=> 1 !!!

rmmod -> "Killed" => need to reboot !!!

83. Fixing kernel crash

```
# ls -l /dev/pcdev- [tab][tab]
```

* 84. Platform device driver matching using platform device ids 254 ~ 256

```
device_id
```

```
static int platform_match(struct device *dev, struct device_driver *drv)    // where matching happens
{
```

```
    struct platform_device *pdev = to_platform_device(dev);
    struct platform_driver *pdrv = to_platform_driver(drv);
```

```
    of_driver_match_device(dev, drv)          // open firmware driver, ie device tree matching
```

```
    acpi_driver_match_device(dev, drv)        //
```

```
    return platform_match_id(pdrv->id_table, pdev); // id_table
```

```
    return (strcmp(pdev->name,drv->name) == 0);
```

```
}
```

```
static const struct platform_device_id *platform_match_id(const struct platform_device_id *id, struct platform_device *pdev)
{
```

```
    while (id->name[0]) {
        if (strcmp(pdev->name, id->name) == 0) {
            pdev->id_entry = id;          <= updated to platform_device->id_entry
            return id;
        }
    }
```

```
    id++;
```

```
}
```

```
return NULL;
```

```
}
```

if init and exit function consists of platform_driver_register and platform_driver_unregister only,

```
=> #define module_platform_driver(__platform_driver) \
    module_driver(__platform_driver, platform_driver_register, \
    platform_driver_unregister)
```

* 85. Fixing error handling in probe function

```
really_probe()
```

-> drv->probe()

7. Device tree

1h26m

258

86. Introduction to device tree

What is device tree?

260 ~ 262

- The "Open Firmware Device Tree", or simply Device Tree (DT), is a data exchange format used for exchanging hardware description data with the software or OS.
- More specifically, it is a description of hardware that is readable by an operating system so that the operating system doesn't need to hard code details of the machine.
- In short, it is a new and recommended way to describe non-discoverable devices(platform devices) to the Linux kernel, which was previously hardcoded into kernel source files.

Source : Documentation/devicetree/usage-model.txt

Device tree

- An operating system uses the Device Tree to discover the topology of the hardware at runtime, and thereby support a majority of available hardware without hardcoded information (assuming drivers were available for all devices)
- The most important thing to understand is that the DT is simply a data structure that describes the hardware. There is nothing magical about it, and it does not magically make all hardware configuration problems go away
- DT provides a language for decoupling the hardware configuration from the device driver and board support files of the Linux kernel (or any other operating system for that matter).
- Using it allows device drivers to become data-driven. To make setup decisions based on data passed into the kernel instead of on permachine hardcoded selections.
- Ideally, a data-driven platform setup should result in less code duplication and make it easier to support a wide range of hardware with a single kernel image.

Why DT is used ?

Linux uses DT for,

- Platform identification : identify the board or machine on which the kernel runs
- Device population:
- The kernel parses the device tree data and generates the required software data structure, which will be used by the kernel code.

Ideally, the device tree is independent of any OS; when you change the OS, you can still use the same device tree file to describe the hardware to the new OS. That is, the device tree makes “adding of device information “ independent of OS

More reading

- https://elinux.org/Device_Tree_What_It_Is
- <https://www.kernel.org/doc/Documentation/devicetree/usage-model.txt>

87. Writing device tree

265

Device tree specification

- You can get the full specification here
- <https://www.devicetree.org/>

Writing device tree

- The device tree supports a hierarchical way of writing hardware description at the soc level, common board level, and board-specific level. Most of the time, writing a new device tree is not difficult, and you can reuse most of the common hardware information from the device tree file of the reference board.
- For example, when you design a new board, which is slightly different from another reference board, then you can reuse the device tree file of the reference board and only add that information which is new in your custom board

Describing hardware hierarchy

- It comes at various level because the board has many device blocks
- SOC
- SOC has an on-chip processor and on-chip peripherals
- The board also has various peripherals onboard, like sensors, LEDs, buttons, joysticks, external memories, touchscreen, etc

Modular approach

mod1 : SOC specific device tree file Board specific device tree file
(This DT file is used as an include file and can be used
with another board which is based on same SOC)

mod2 : Board specific device tree file

linux/arch/arm/boot/dts/am335x-evm.dts => #include "am33xx.dtsi" (soc level device file)

```
am335x-boneblack.dts                                272, 278
#include "am33xx.dtsi"
#include "am335x-bone-common.dtsi"
#include "am335x-boneblack-common.dtsi"
```

88. Device tree structure 19

Overview of device tree structure

- ✓ Device tree is a collection of device nodes
- ✓ A 'device node' or simply called 'a node' represents a device.
Nodes are organized in some systematic way inside the device tree file.
- ✓ They also have parent and child relationship, and every device tree must have one root node
- ✓ A node explains itself, that is, reveals its data and resources using its "properties."

Root node

- The device tree has a single root node of which all other device nodes are descendants.
The full path to the root node is /.
- All device trees shall have a root node, and the following nodes shall be present at the root of all device trees:
 - One /CPUs node
 - At least one /memory node

Chapter 3 :DEVICE NODE REQUIREMENTS Devicetree Specification Release v0.3a

How to write a device tree ?

- Remember that you most probably be writing device tree addons or overlays for your board-related changes but not for entire soc.
- The soc specific device tree will be given by the vendor in the form of device tree inclusion file (.dtsi) and you just need to include that in your board-level device tree
- Follow modulatory approach while writing device tree

```
am335x-boneblack.dts                                272, 278
```

```
/ {
    model = "TI AM335x BeagleBone Black";           <= added
    compatible = "ti, am335x-bone-black", ... ;    <= override that of included files
```

```

    };                                <= root level node (children of root)
};                                    <= root

    &sgx {                             <= reference
        status = "okay";
    };

am33xx.dtsi

    / {
        compatible = "ti, am33xx";      <= overridden

        sgx: sgx@56000000 {
            compatible = " ";
            ti, hwmods = "gfx";
        };
    };
};

#include "am335x-bone-common.dtsi"
#include "am335x-boneblack-common.dtsi"

```

Never edit top level files => it will be overridden later.

89. Device tree writing syntax

20

Device tree writing syntax

279

- Node name
- Node Label
- Standard and non-standard property names
- Different data type representation (u32, byte, byte stream, string, stream of strings, Boolean, etc)
- SoC node and children

Node name

- Refer device tree specification Release v0.3 from devicetree.org

<https://github.com/devicetree-org/devicetree-specification/releases/tag/v0.4>

The unit-address component of the name is specific to the bus type on which the node sits. It consists of one or more ASCII characters from the set of characters in Table 2.1. The unit-address must match the first address specified in the reg property of the node.

If the node has no reg property, the @unit-address must be omitted and the node-name alone differentiates the node from other nodes at the same level in the tree.

The binding for a particular bus may specify additional, more specific requirements for the format of reg and the unit-address.

ex.

281

```
i2c0: i2c@44e0b000 {
    compatible = "ti, omap4-i2c";
    #address-cells = <1>;
    #size-cells = <0>;
    ti, hwmods = "i2c1";
    reg = <0x44e0b000 0x1000>;      <= AM335x TRM ch2 memory map
    interrupts = <70>;
    status = "disabled";
};

i2c1: i2c@44e2a000 {
    compatible = "ti, omap4-i2c";
    #address-cells = <1>;
    #size-cells = <0>;
    ti, hwmods = "i2c2";
    reg = <0x4802a000 0x1000>;
    interrupts = <71>;
    status = "disabled";
};

i2c2: i2c@48060000 {
    ...
}
```

90. Device tree parent and child node

282 ~ 283

```
&i2c0 {
    status = "okay";

    tps: tps@24 {
        reg = <0x24>;    // i2c address
    };

    baseboard_eeprom: baseboard_eeprom@50 {
        reg = <0x50>

        baseboard_data: baseboard_data@0 {
            reg = <0 0x100>;
        };
    };
};
```



```
label,
gpios,                : standard property
default-state

linux,default-trigger : custom property
```

Different types of properties

- Standard property
- Custom property (non-standard)
- Standard properties are those which is explained by the specification and the device-driver binding documentation
- Custom properties are specific to a particular vendor or organization which is not documented by the specification.
- That is why, when you use custom property, always begin with your organization name.

Root compatible property of BBB

```
/ {
    model = "TI AM335x BeagleBone Black";
    compatible = "ti,am335x-bone-black", "ti,am335x-bone", "ti,am33xx";
};

    Sorted string list from most compatible to least.
```

; Root compatible property is used for machine identification

Uses of compatible property

1. Machine identification and initialization
2. Match and load the appropriate driver for the device

```
&i2c0 {
    status = "okay";                !!!!

    tps: tps@24 {
        reg = <0x24>;    // i2c address
    };

    baseboard_eeprom: baseboard_eeprom@50 {
        compatible = "atmel,24c256";    <= match found in linux/drivers/misc/eeprom/at24.c
        reg = <0x50>

        #address-cells = <1>;
        #size-cells = <1>;
```

```

        baseboard_data: baseboard_data@0 {
            reg = <0 0x100>;
        };
};

i2c0: i2c@44e0b000 {
    compatible = "ti, omap4-i2c";           <= matched !
    #address-cells = <1>;
    #size-cells = <0>;
    ti, hwmods = "i2c1";
    reg = <0x44e0b000 0x1000>;
    interrupts = <70>;
    status = "disabled";
};

```

in linux/drivers/i2c/buses/i2c-omap.c -> drivers

```

static struct platform_driver omap_i2c_driver = {
    .probe = omap_i2c_probe,
    .remove = omap_i2c_remove,
    .driver = {
        .name = "omap_i2c",
        .pm = &omap_i2c_pm_ops,
        .of_match_table = of_match_ptr(omap_i2c_of_match),
    },
}

static const struct of_device_id omap_i2c_of_match[] = {
    {
        .compatible = "ti, omap4-i2c",           <= matched !
        .data = &omap4_pdata,
    },
    {
        .compatible = "ti, omap3-i2c",
        .data = &omap3_pdata,
    },
    ...
}

```

Device tree bindings

- How do you know which property name and value pair should be used to describe a node in the device tree?
- Device tree binding document. The driver writer must document these details
- The properties that are necessary to describe a device in the device tree depends on the requirements of the Linux driver

for that device

- When all the required properties are provided, the driver of charge can detect the device from the device tree and configure it.

linux/Documentation/devicetree/bindings/i2c/i2c-omap.txt

Required properties:

- compatible : Must be

"ti,omap2420-i2c" for OMAP2420 SoCs

"ti,omap2430-i2c" for OMAP2430 SoCs

"ti,omap3-i2c" for OMAP3 SoCs

"ti,omap4-i2c" for OMAP4+ SoCs

"ti,omap654-i2c", "ti,omap4-i2c" for AM654 SoCs

"ti,j721e-i2c", "ti,omap4-i2c" for J721E SoCs

Recommended properties:

Optional properties:

ex...

- Compatible strings and properties are first defined by the client program (OS , drivers) then shared with DT writer

Device tree bindings- points to remember

- Case 1 : When the driver is already available in the Linux kernel for the device 'x,' but you just need to add device 'x' entry in the device tree then you must consult 'x' drivers binding document which guides you through creating device tree node for device 'x.'
- Case 2 : When the driver is not available for the device 'x,' then you should write you own driver, you should decide what properties to use (could be a combination of standard and non-standard property), you should then provide the device tree binding document describing what are all the properties and compatible strings a device tree write must include.

ex. lm75

linux/Documentation/devicetree/bindings/hwmon/lm75.txt

```
sensor@48 {  
    compatible = "st,stlm75";  
    reg = <0x48>;  
}
```

linux/drivers/hwmon/lm75.c

```
static const struct of_device_id lm75_of_match[] = {
    .compatible = "ti, tmpXXX";
    .data = (void *)tmpXXX
};
```

ex. mpu 6050

[linux/Documentation/devicetree/bindings/iio/imu/inv_mpu6050.txt](#)

```
mpu6050@68 {
    compatible = "invensense,mpu6050";
    reg = <0x68>;
    ...
}
```

[linux/drivers/hwmon/lm75.c](#)

```
static const struct of_device_id lm75_of_match[] = {
    .compatible = "ti, tmpXXX";
    .data = (void *)tmpXXX
};
```

<https://kernel.org/doc/Documentation/devicetree/bindings/i2c/i2c-omap.txt>

Linux conventions to write device tree

hex constants are lower case

- use "0x" instead of "0X"
- use a..f instead of A..F, eg 0xf instead of 0xF

node names

- should begin with a character in the range 'a' to 'z', 'A' to 'Z'
- unit-address does not have a leading "0x" (the number is assumed to be hexadecimal)
- unit-address does not have leading zeros
- use dash "-" instead of underscore "_"

label names

- should begin with a character in the range 'a' to 'z', 'A' to 'Z'
- should be lowercase
- use underscore "_" instead of dash "-"

property names

- should be lower case
- should begin with a character in the range 'a' to 'z'
- use dash "-" instead of underscore "_"

https://elinux.org/Device_Tree_Linux#Linux_vs_ePAPR_Version_1.1

94. pcd device tree version

```
# git branch -a
*4.14
  remote/origin/5.4
  remote/origin/HEAD -> origin/4.14

# git stash && git checkout 5.4

# git branch
  4.14
  4.9
  *5.4

# git stash apply

# git clone https://github.com/beagleboard/linux.git linux_bbb_5.4
```

95. Switching to Linux kernel version 5.4

```
# git checkout 5.4
```

96. Updating Linux kernel image of 5.4

```
make ARCH=arm CROSS_COMPILE=arm-linux-gnueabihf- bb.org_defconfig
make ARCH=arm CROSS_COMPILE=arm-linux-gnueabihf- uImage dtbs LOADADDR=0x80008000 -j4
make ARCH=arm CROSS_COMPILE=arm-linux-gnueabihf- -j4 modules
make ARCH=arm CROSS_COMPILE=arm-linux-gnueabihf- modules_install
```

```
$ cd /home/kiran/workspace/ldd/src/linux_bbb_5.4
$ cd arch/arm/boot
$ cp uImage /media/kiran/BOOT/
$ cd dtb
$ cp am335x-boneblack.dtb /media/kiran/BOOT/
$ cd /lib/modules/
$ sudo cp -a 5.4.47/ /media/kiran/ROOTFS/lib/modules/
$ sync
```

```
$ uname -r
5.4.47
```

```
$ ifconfig -> reboot if usb0 & usb1 is not seen.
```

8. device tree nodes and platform driver 1h16m

97. Device tree nodes for pcd driver

```
workspace/ldd/src/linux_bbb_5.4/arch/arm/boot/dts# vi am335x-boneblack.dts
#include "am335x-boneblack-lddcrs.dtsi"

workspace/ldd/src/linux_bbb_5.4/arch/arm/boot/dts# vi am335x-boneblack-lddcourse.dtsi
/ {
    pcdev-1 {
        org,size = <512>
        org,device-serial-num = "PCDEV1ABC123"
        org,perm = <0x11>
    };
    pcdev-2 {
        org,size = <1024>
        org,device-serial-num = "PCDEV2ABC123"
        org,perm = <0x11>
    };
    pcdev-3 {
        org,size = <256>
        org,device-serial-num = "PCDEV3ABC123"
        org,perm = <0x11>
    };
    pcdev-1 {
        org,size = <1024>
        org,device-serial-num = "PCDEV4ABC123"
        org,perm = <0x11>
    };
};
```

-> device tree compiler to get dtb

98. Pcd platform driver DT coding part-1

Platform_device_id can be used only when device register function is called manually.
=> another list of device is needed inside device tree.

```
static int platform_match(struct device *dev, struct device_driver *drv)
{
    of_driver_match_device(dev, drv)                // open firmware driver, ie device tree matching

    // Refer to 84 for details.

}
```

```

static inline int of_of_driver_match_device(struct device *dev, const struct device_driver *drv)
{
    return of_match_device(drv->of_match_table, dev) != NULL;
    ^^^^^^^^^^^^^^^^^^^
}

struct of_device_id {
    char name[32];
    char type[32];
    char compatible[128];
    const void *data;
};

const struct of_device_id *of_match_device(const struct of_device_id *matches, const struct device *dev)
{
    if ((!matches) || (!dev->of_node)) return NULL;
    return of_match_node(matches, dev->of_node);
}

struct device_node *of_node;
struct device_node {
    const char *name;

    struct property *properties;
    struct device_node *parent;
    struct device_node *child;
    struct device_node *sibling;
    ...
};

```

99. Testing device tree changes on board

>> on VBox

```

/ {
    pcdev1: pcdev-1 {
        compatible = "pcdev-E1x","pcdev-A1x";
        org,size = <512>;
        org,device-serial-num = "PCDEV1ABC123";
        org,perm = <0x11>;
    };
};

# make ARCH=arm CROSS_COMPILE=arm-linux-gnueabi- am335x-boneblack.dtb
# scp arch/arm/boot/dts/am335x-boneblack.dtb debian@192.168.7.2:/home/debian/drivers

or in Makefile,

```

```

copy-dtb:
  scp ~/workspace/ldd/source/linux_bbb_5.4/arch/arm/boot/dts/am335x-boneblack.dtb
  debian@192.168.7.2:/home/debian/drivers

copy-drv:
  scp *.ko  debian@192.168.7.2:/home/debian/drivers

```

>> on beaglebone

```

# mount /dev/mmcblk0p1 /mnt/
# cd /mnt/
# ls
  am335x-boneblack_4.14.dtb      MLO      uImage      uEnv.txt
# cp /home/debian/drivers/am335x-boneblack.dtb .
# sync
# reboot

# ls /sys/devices/platform/
...      pcdev-1      pcdev-2      pcdev-3      pcdev-4      ...
# ls /sys/devices/platform/pcdev-1
driver_override      modalias      of_node      power      subsystem      uevent
# ls /sys/devices/platform/pcdev-1/of_node
compatible      name      org,device-serial-num      org,perm      org,size
# cat /sys/devices/platform/pcdev-1/of_node/compatible
pcdev-E1xpcdev-A1x

```

```
$ sudo insmod pcd_platform_driver_dt.ko
```

```

$ dmesg
[ 168.xx] pcd_platform_driver_dt: loading out-of-tree module taint kernel.
[ 168.xx] pcd_platform_driver_probe : A device is detected
[ 168.xx] pcd_platform_driver_probe : A device is detected
[ 168.xx] pcd_platform_driver_probe : A device is detected
[ 168.xx] pcd_platform_driver_probe : A device is detected
[ 168.xx] pcd_platform_driver_init : pcd platform driver loaded

```

```

$ sudo rmmod pcd_platform_driver_dt
[ 168.xx] pcd_platform_driver_remove : A device is remove
[ 168.xx] pcd_platform_driver_remove : A device is remove
[ 168.xx] pcd_platform_driver_remove : A device is remove
[ 168.xx] pcd_platform_driver_remove : A device is remove
[ 168.xx] pcd_platform_driver_cleanup : pcd platform driver unloaded

```

```

struct platform_device {
    struct device      dev;

```



```

{
    struct device_node      *of_node;    // represents an associated device tree node
    struct fwnode_handle    *fwnode;    // fw device node
};
};

```

linux/include/of.h

```

of_property_read_string(const struct device_node *np,
                        const char *propname,
                        const char **out_string);

```

pr_info(...) vs dev_info(dev, ...) -> can't use in _init & _exit where dev is not available)

101. Pcd platform driver DT coding part-3

102. Pcd platform driver DT coding part-4

9 device tree overlays	54m	317
------------------------	-----	-----

- 103. Introduction to device tree overlays
- 104. Device tree overlays exercise
- 105. Updating u-boot
- 106. Updating u-boot contd.
- 107. Applying overlays using u-boot commands
- 108. Applying overlays and testing using u-boot uEnv.txt file

10 Linux device driver model	1h32m	341
------------------------------	-------	-----

- 109. Linux device model
- 110. kobjects
- 111. kobject type and kset
- 112. sysfs and kobject attributes
- 113. Creating sysfs attributes
- 114. pcd_sysfs_attributes coding part 1

- 115. pcd_sysfs_attributes coding part 2
- 116. pcd_sysfs kernel module testing
- 117. show and store methods of the sysfs attributes
- 118. pcd_sysfs_attributes coding part 3
- 119. Attribute grouping

11 Linux GPIO subsystem 32m ???

- 120. Introduction
- 121. GPIOs of BBB
- 122. Pad configuration register
- 123. Linux GPIO subsystem
- 124. Consumer accessing GPIO pins

12 GPIO sysfs driver implementation 2h

- 125. Exercise GPIO Sysfs driver implementation
- 126. GPIO Sysfs driver implementation part-1
- 127. GPIO Sysfs driver implementation part-2
- 128. GPIO Sysfs driver implementation part-3
- 129. GPIO Sysfs driver implementation part-4
- 130. GPIO Sysfs driver implementation part-5
- 131. GPIO Sysfs driver implementation part-6
- 132. GPIO Sysfs driver implementation part-7
- 133. GPIO Sysfs driver implementation part-8
- 134. Gpio syfs driver testing

13 pin control subsystem of Linux 2h2m

- 135. Pin control subsystem and pin controller
- 136. Writing pin configuration node
- 137. LCD exercise
- 138. Significance of LCD application
- 139. Adding pin configuration node for gpios
- 140. Exploring LCD code
- 141. Sending command to LCD
- 142. Creating LCD command code
- 143. Testing LCD application over gpio sysfs
- 144. Assignment : Implementing LCD platform driver

14 Linux synchronization services

1h35m

- 145. Avoiding race conditions
- 146. Linux locking services
- 147. Spinlock Vs Mutex
- 148. Linux spinlock functions
- 149. Mutex
- 150. Using locking functions in the code
- 151. BONUS LECTURE